

Documentation_Capstone_Propensity

January 3, 2025

This project focuses on developing a propensity model for an insurance company to predict which potential customers are likely to respond positively to marketing campaigns. Using historical customer data (`train.csv`) and a list of potential leads (`test.csv`), the model will forecast customer engagement, helping optimize marketing strategies and resource allocation.

1 Step 1: Import libraries

```
[2]: # Importing libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import random

from sklearn.preprocessing import power_transform
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.linear_model import RidgeClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.ensemble import RandomForestClassifier
from imblearn.over_sampling import SMOTE
from sklearn.preprocessing import StandardScaler, PowerTransformer
from sklearn.preprocessing import OneHotEncoder
from sklearn.decomposition import PCA
from sklearn.metrics import accuracy_score, confusion_matrix, \
    roc_auc_score, classification_report
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import FunctionTransformer
from sklearn.model_selection import GridSearchCV
from scipy.stats import chi2_contingency
from scipy import stats
import joblib

# suppress warnings
import warnings
warnings.filterwarnings('ignore')
```

2 Step 2: Load dataset

```
[3]: df = pd.read_excel('/content/drive/MyDrive/Upgrad/Data sets/Capstone/train.  
↳xlsx')
```

3 Step 3: Exploratory Data Analysis

3.1 Understand the Basic Structure

```
[4]: # Reading the data with all the columns visible  
pd.options.display.max_columns=None  
df.head(5)
```

```
[4]:
```

	custAge	profession	marital	schooling	default	housing	loan	\
0	34.0	admin.	single	university.degree	no	no	yes	
1	31.0	services	single	high.school	no	no	no	
2	NaN	admin.	single	high.school	no	no	no	
3	52.0	admin.	divorced	university.degree	unknown	yes	no	
4	39.0	blue-collar	single	NaN	unknown	yes	no	

	contact	month	day_of_week	campaign	pdays	previous	poutcome	\
0	cellular	apr	wed	2.0	999.0	0.0	nonexistent	
1	cellular	jul	thu	35.0	999.0	0.0	nonexistent	
2	telephone	jun	NaN	1.0	999.0	0.0	nonexistent	
3	cellular	jul	tue	2.0	999.0	0.0	nonexistent	
4	cellular	jul	tue	6.0	999.0	0.0	nonexistent	

	emp.var.rate	cons.price.idx	cons.conf.idx	euribor3m	nr.employed	\
0	-1.8	93.075	-47.1	1.498	5099.1	
1	1.4	93.918	-42.7	4.968	5228.1	
2	1.4	94.465	-41.8	4.961	5228.1	
3	1.4	93.918	-42.7	4.962	5228.1	
4	1.4	93.918	-42.7	4.961	5228.1	

	pmonths	pastEmail	responded	profit	id
0	999.0	0.0	no	NaN	1.0
1	999.0	0.0	no	NaN	2.0
2	999.0	0.0	no	NaN	3.0
3	999.0	0.0	no	NaN	4.0
4	999.0	0.0	no	NaN	5.0

```
[5]: df.columns
```

```
[5]: Index(['custAge', 'profession', 'marital', 'schooling', 'default', 'housing',  
        'loan', 'contact', 'month', 'day_of_week', 'campaign', 'pdays',  
        'previous', 'poutcome', 'emp.var.rate', 'cons.price.idx',  
        'cons.conf.idx', 'euribor3m', 'nr.employed', 'pmonths', 'pastEmail',
```

```
    'responded', 'profit', 'id'],
    dtype='object')
```

```
[6]: # Drop unwanted features based on image
df = df.drop(columns=['profit', 'id'], axis = 1)
```

```
[7]: # Get the rows and columns of training data
df_shape = df.shape
print("Data shape:", df_shape)
```

Data shape: (8240, 22)

```
[8]: # Get basic information about data types and non-null values
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8240 entries, 0 to 8239
Data columns (total 22 columns):
#   Column                Non-Null Count  Dtype
---  -
0   custAge                6224 non-null   float64
1   profession              8238 non-null   object
2   marital                8238 non-null   object
3   schooling               5832 non-null   object
4   default                 8238 non-null   object
5   housing                 8238 non-null   object
6   loan                    8238 non-null   object
7   contact                 8238 non-null   object
8   month                   8238 non-null   object
9   day_of_week             7451 non-null   object
10  campaign                8238 non-null   float64
11  pdays                  8238 non-null   float64
12  previous                8238 non-null   float64
13  poutcome                8238 non-null   object
14  emp.var.rate            8238 non-null   float64
15  cons.price.idx          8238 non-null   float64
16  cons.conf.idx           8238 non-null   float64
17  euribor3m               8238 non-null   float64
18  nr.employed             8238 non-null   float64
19  pmonths                 8238 non-null   float64
20  pastEmail               8238 non-null   float64
21  responded               8238 non-null   object
dtypes: float64(11), object(11)
memory usage: 1.4+ MB
```

3.2 Summarize the Data

```
[9]: # Statistic description of numerical columns
df.describe()
```

```
[9]:
```

	custAge	campaign	pdays	previous	emp.var.rate \
count	6224.000000	8238.000000	8238.000000	8238.000000	8238.000000
mean	39.953728	2.531682	960.916606	0.183054	0.056397
std	10.540516	2.709773	190.695054	0.514209	1.566550
min	18.000000	1.000000	0.000000	0.000000	-3.400000
25%	32.000000	1.000000	999.000000	0.000000	-1.800000
50%	38.000000	2.000000	999.000000	0.000000	1.100000
75%	47.000000	3.000000	999.000000	0.000000	1.400000
max	94.000000	40.000000	999.000000	6.000000	1.400000

	cons.price.idx	cons.conf.idx	euribor3m	nr.employed	pmonths \
count	8238.000000	8238.000000	8238.000000	8238.000000	8238.000000
mean	93.570977	-40.577907	3.586929	5165.575965	960.687436
std	0.578782	4.650101	1.742784	72.727423	191.841012
min	92.201000	-50.800000	0.634000	4963.600000	0.000000
25%	93.075000	-42.700000	1.334000	5099.100000	999.000000
50%	93.444000	-41.800000	4.857000	5191.000000	999.000000
75%	93.994000	-36.400000	4.961000	5228.100000	999.000000
max	94.767000	-26.900000	5.045000	5228.100000	999.000000

	pastEmail
count	8238.000000
mean	0.365501
std	1.294101
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	25.000000

- The maximum values observed in the 'pdays' and 'pmonths' columns appear to be the result of missing data.
- These values will be appropriately encoded in the Data Cleaning process.

3.3 Check for Duplicate values

```
[10]: # Check for duplicates
df.duplicated().sum()
```

```
[10]: 37
```

```
[11]: # Data have 37 duplicate records
print('Original Shape of Data: ',df.shape)

# Remove duplicates and resetting the index
```

```
df = df.drop_duplicates().reset_index(drop=True)
print('Shape of Data after removing duplicates: ',df.shape)
```

Original Shape of Data: (8240, 22)

Shape of Data after removing duplicates: (8203, 22)

3.4 Check for Missing values

```
[12]: # Checking null value count for each column
df.isnull().sum()
```

```
[12]: custAge      2000
      profession      1
      marital        1
      schooling    2394
      default        1
      housing        1
      loan           1
      contact        1
      month          1
      day_of_week   785
      campaign       1
      pdays          1
      previous       1
      poutcome       1
      emp.var.rate   1
      cons.price.idx  1
      cons.conf.idx  1
      euribor3m      1
      nr.employed    1
      pmonths        1
      pastEmail      1
      responded      1
      dtype: int64
```

```
[13]: # Filtering the rows with null values in target column ('responded')
df[df['responded'].isnull()]
```

```
[13]:      custAge profession marital schooling default housing loan contact month \
8202      NaN      NaN      NaN      NaN      NaN      NaN  NaN  NaN      NaN  NaN

      day_of_week campaign pdays previous poutcome emp.var.rate \
8202      NaN      NaN      NaN      NaN      NaN      NaN      NaN

      cons.price.idx cons.conf.idx euribor3m nr.employed pmonths \
8202      NaN      NaN      NaN      NaN      NaN      NaN

      pastEmail responded
```

8202 NaN NaN

- It was identified that a record contains null values in all columns including target column, hence this record will be removed.

```
[14]: # Removing the record containing null value in target column
df = df.dropna(subset=['responded'])
```

```
[15]: # Checking for null percentage for each column
round((df.isnull().sum() / len(df)) * 100,2)
```

```
[15]: custAge          24.37
      profession      0.00
      marital         0.00
      schooling       29.18
      default         0.00
      housing         0.00
      loan            0.00
      contact         0.00
      month           0.00
      day_of_week     9.56
      campaign        0.00
      pdays           0.00
      previous        0.00
      poutcome        0.00
      emp.var.rate    0.00
      cons.price.idx  0.00
      cons.conf.idx   0.00
      euribor3m       0.00
      nr.employed     0.00
      pmonths         0.00
      pastEmail       0.00
      responded       0.00
      dtype: float64
```

3.5 Target Variable Analysis


```
[16]: # Calculate value counts of Target column
df['responded'].value_counts()
```

```
[16]: responded
no      7274
yes      928
Name: count, dtype: int64
```

```
[17]: # Calculate percentage distribution of values in target column
round(df['responded'].value_counts(normalize=True) * 100,2)
```

```
[17]: responded
      no      88.69
      yes     11.31
      Name: proportion, dtype: float64
```

```
[18]: sns.countplot(df, x='responded')
      plt.title("Target Variable")
      plt.show()
```



The target variable in the dataset is heavily skewed, as around 88% of the customers did not engage with the marketing campaign, while only 11% responded. To ensure accurate model performance, it is vital to address this class imbalance prior to model development.

3.6 Numerical Feature Analysis

```
[19]: num_cols = df._get_numeric_data().columns
      num_cols
```


```
[19]: Index(['custAge', 'campaign', 'pdays', 'previous', 'emp.var.rate',
          'cons.price.idx', 'cons.conf.idx', 'euribor3m', 'nr.employed',
          'pmonths', 'pastEmail'],
```

```
dtype='object')
```


```
[20]: # Boxplot and Histogram for numerical columns
```

```
for col in num_cols:
    fig, axes = plt.subplots(nrows=1, ncols=3, figsize=(15,3))
    skewness = round(df[col].skew(),2)
    axes[0].text(0.5, 0.5, (f"{col}\nSkewness: {skewness}"), fontsize=12,
    ↪ha='center', va='center')
    axes[0].axis('off')
    axes[1].boxplot(df[col].dropna())
    axes[2].hist(df[col])
    # axes[0].set_xlabel(col)
plt.show()
```


custAge
Skewness: 0.86




campaign
Skewness: 4.82




pdays
Skewness: -4.8




previous
Skewness: 3.83




emp.var.rate
Skewness: -0.67




cons.price.idx
Skewness: -0.19




cons.conf.idx
Skewness: 0.3




euribor3m
Skewness: -0.66




nr.employed
Skewness: -1.02



pmonths
Skewness: -4.8



pastEmail
Skewness: 6.04



- From the above charts we can observe that custAge, campaign, previous and pastEmail columns are right skewed.
- nr.employed column is left skewed.
- pdays and pmonths column are to be treated after handling the missing values.

3.7 Categorical Feature - Univariate Analysis

```
[21]: cat_cols = df.drop(columns=num_cols, axis=1).columns
      cat_cols
```

```
[21]: Index(['profession', 'marital', 'schooling', 'default', 'housing', 'loan',
           'contact', 'month', 'day_of_week', 'poutcome', 'responded'],
          dtype='object')
```

```
[22]: for column in cat_cols:
      print(df[column].value_counts())
      print('\n')
```

```
profession
admin.          2090
blue-collar     1842
technician     1340
services        790
management      580
retired         335
entrepreneur    314
self-employed   279
housemaid       213
unemployed      189
student         159
unknown         71
Name: count, dtype: int64
```

```
marital
married      4933
single       2329
divorced      930
unknown       10
Name: count, dtype: int64
```

```
schooling
university.degree  1716
high.school        1337
basic.9y           862
professional.course  736
basic.4y           585
basic.6y           312
unknown            260
illiterate          1
Name: count, dtype: int64
```

```
default
no          6587
unknown     1614
yes         1
Name: count, dtype: int64
```

```
housing
yes        4281
no         3737
unknown    184
Name: count, dtype: int64
```

```
loan
no         6740
yes        1278
unknown    184
Name: count, dtype: int64
```

```
contact
cellular   5211
telephone  2991
Name: count, dtype: int64
```

```
month
may        2809
jul        1344
aug        1225
jun        1054
nov         808
apr         551
oct         156
sep         120
mar         106
dec         29
Name: count, dtype: int64
```

```
day_of_week
mon        1590
thu        1525
tue        1473
wed        1468
fri        1362
Name: count, dtype: int64
```

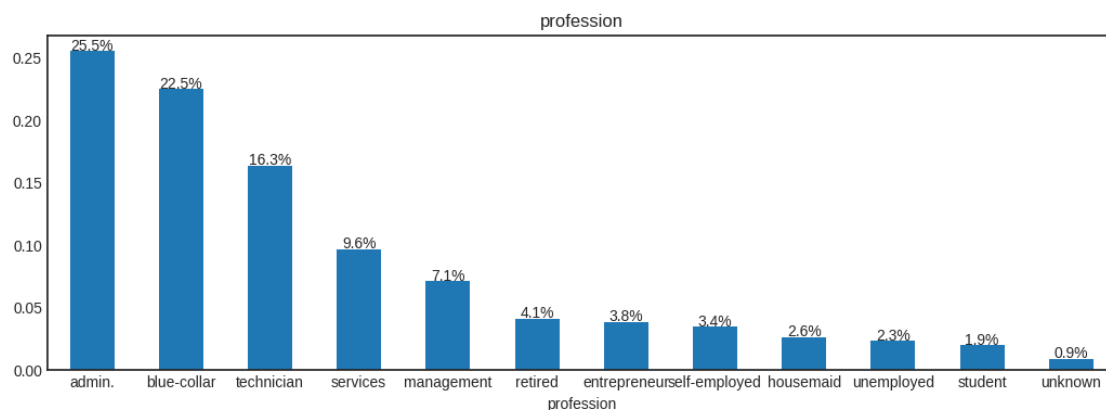
```
poutcome
nonexistent    7025
failure        894
success        283
Name: count, dtype: int64
```


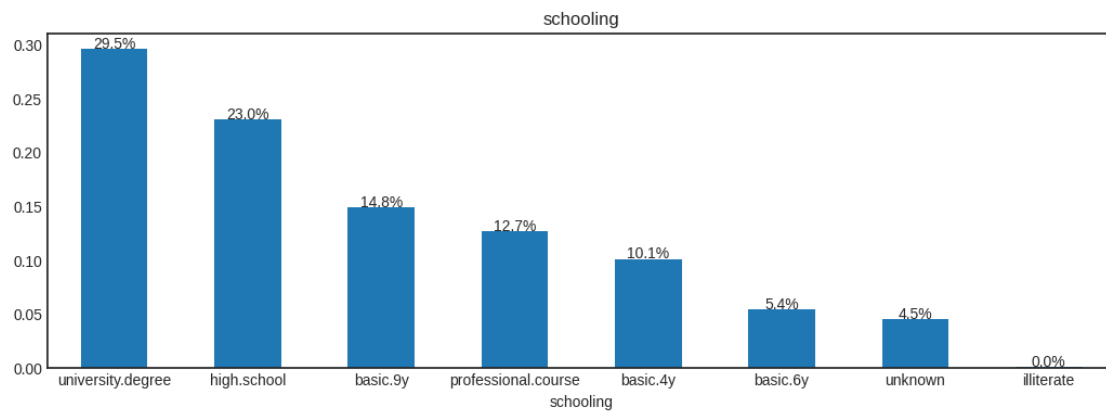
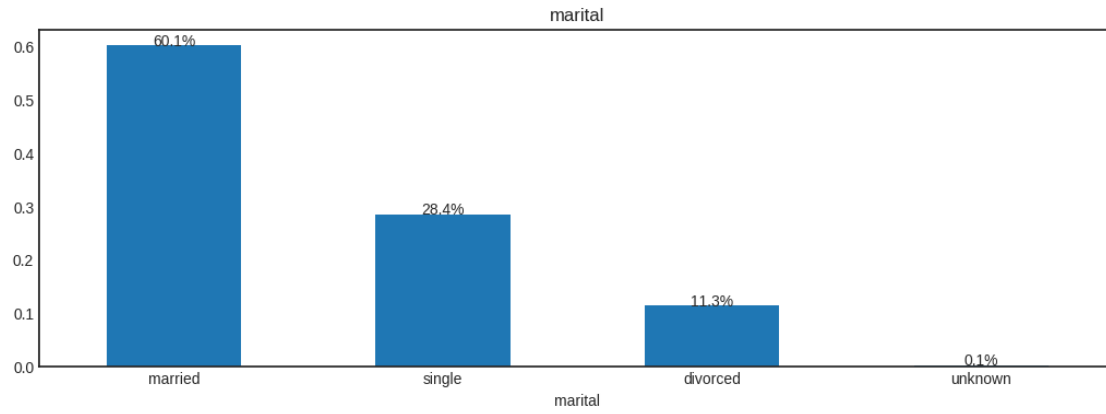
```
responded
no      7274
yes      928
Name: count, dtype: int64
```


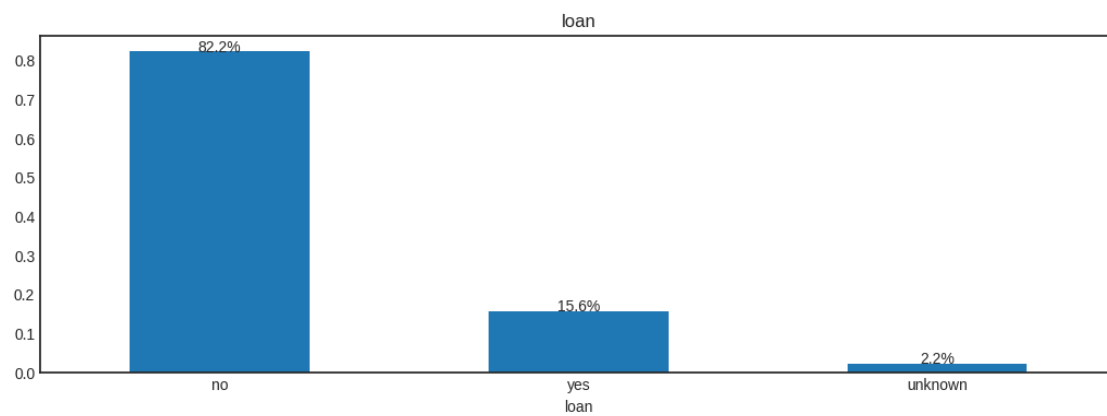

```
[27]: # plotting bar chart for each categorical variable
plt.style.use('seaborn-v0_8-white')

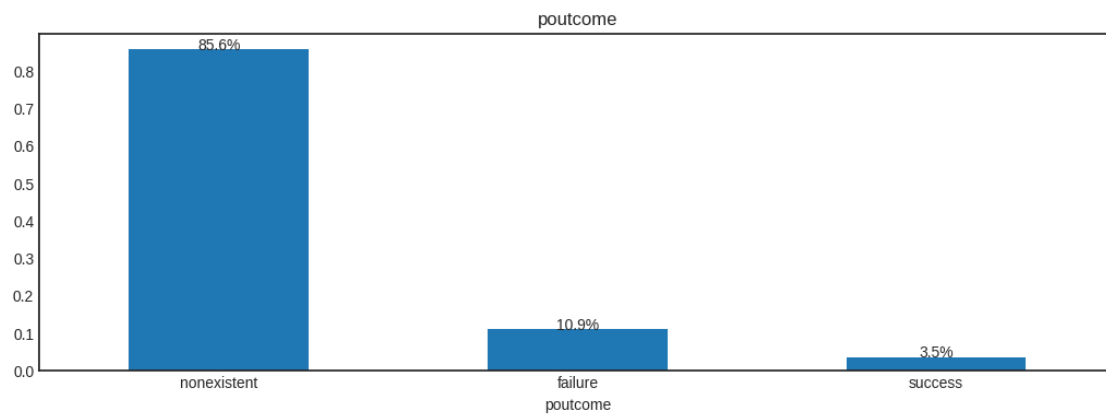
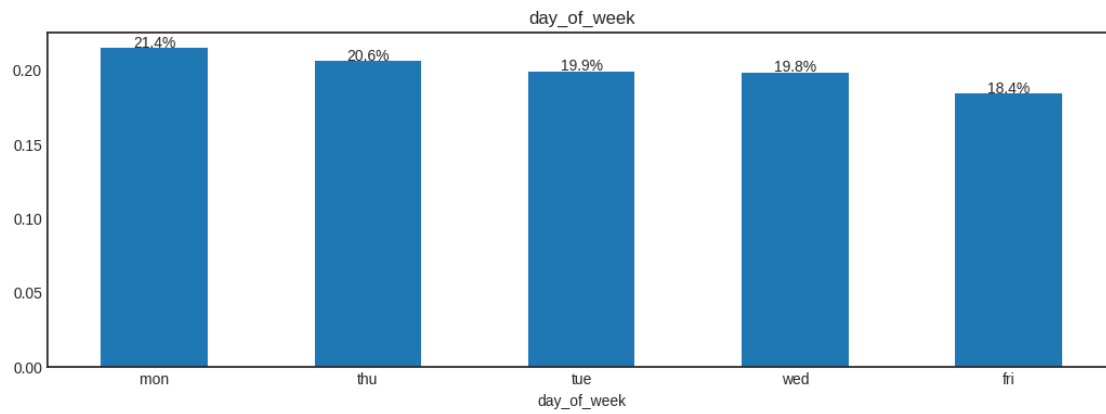

for column in cat_cols:
    plt.figure(figsize=(28,4))
    ax = plt.subplot(121)
    df[column].value_counts(normalize=True).plot(kind="bar")
    plt.xticks(rotation= 0)
    plt.title(column)


    # Add percentage labels to the top of each bar
    for p in ax.patches:
        ax.text(p.get_x()+p.get_width()/2., p.get_height()+0.001, f"{p.
↪get_height()*100:.1f}%", ha="center")
```











4 Step 4: Data Cleaning

4.1 Data Aggregation

Schooling

- basic.4y, basic.6y, basic.9y education can be grouped as Primary Education
- Since illiterate has only one record, grouping it into unknown section

```
[376]: df['schooling'] = df['schooling'].replace(['basic.4y', 'basic.6y', 'basic.
↪9y'], 'primary.education')
df['schooling'] = df['schooling'].replace('illiterate', 'unknown')
```

```
[377]: df['schooling'].value_counts()
```

```
[377]: schooling
primary.education    1759
university.degree    1716
high.school          1337
professional.course    736
unknown              261
Name: count, dtype: int64
```

4.2 Handle Missing Data

‘custAge’, ‘schooling’ have 25% of missing data and ‘day of the week’ has around 9% of missing data.

- Customer age can affect responses to insurance marketing based on different life stages
- Day of the Week affects availability for making decisions
- Schooling reflects educational background, which may impact the likelihood of purchasing insurance.

Dropping these variables would result in a significant loss of information. Therefore, we will use different imputation methods to fill the missing values in these columns.

4.2.1 Schooling

The Schooling column may have an impact on an individual's profession, as education level often correlates with career choices. To address the missing values in the Schooling column, we will analyze the relationship between Schooling and Profession to identify patterns and use this relationship to impute the missing data effectively.

```
[378]: # Create a cross-tab for 'schooling' and 'profession'
cross_tab = pd.crosstab(df['schooling'], df['profession'], normalize =_
    ↪ 'index')*100

# Set up the matplotlib figure
plt.figure(figsize=(20, 5))

# Create a heatmap for the cross-tabulation
sns.heatmap(cross_tab, annot=True, fmt=".2f", cmap='Blues', cbar=True, _
    ↪ linewidths=0.5)
plt.gca().xaxis.set_label_position('top') # Moves the xlabel to the top
plt.gca().xaxis.set_ticks_position('top') # Moves the xticks to the top
plt.show()
```



Based on the analysis of the Schooling and Profession columns, we observe distinct patterns linking education levels to specific professions. To handle missing values in the Schooling column, we will impute them by associating education levels with their corresponding professions. Any remaining missing data will be categorized as “Unknown.”

```
[379]: schooling_profession_mapping = {
    'technician': 'professional.course',
    'blue-collar': 'primary.education',
    'admin.': 'university.degree',
    'services': 'high.school'
}

# Function to impute missing 'Schooling' values based on 'Profession'
def impute_schooling(row):
    if pd.isnull(row['schooling']):
```

```

        return schooling_profession_mapping.get(row['profession'], 'unknown')
    else:
        return row['schooling']

# Apply the function to impute missing values in 'Schooling'
df['schooling'] = df.apply(impute_schooling, axis=1)


```

4.2.2 Customer Age

```

[380]: plt.figure(figsize=(15,5))
sns.boxplot(x='profession', y='custAge', data=df)
plt.show()

```



To address the missing values in the Age column, an analysis was performed to understand the relationship between Profession and Age. The analysis revealed that retired individuals have a higher average age, while students have a lower average age compared to other professions. Based on this, missing Age values will be imputed by using the mean age specific to the retired and student professions.

```

[381]: # Calculate mean age for retired, student, and other professions
mean_ages = {
    'retired': df[df['profession'] == 'retired']['custAge'].mean(),
    'student': df[df['profession'] == 'student']['custAge'].mean(),
    'other': df[~df['profession'].isin(['retired', 'student'])]['custAge'].
    ↪mean()
}

# Use vectorized operations to fill missing custAge values based on profession
df['custAge'] = df.apply(
    lambda row: mean_ages['retired'] if row['profession'] == 'retired' and pd.
    ↪isnull(row['custAge']) else
    mean_ages['student'] if row['profession'] == 'student' and pd.
    ↪isnull(row['custAge']) else

```

```

        mean_ages['other'] if pd.isnull(row['custAge']) else
    row['custAge'],
    axis=1
)


```

```

[382]: df['custAge'].hist()
skewness = round(df['custAge'].skew(),2)
print(f"Skewness: {skewness}")

```

Skewness: 1.01



4.2.3 Day of the week

There is no clear relationship observed between the 'day_of_week' column and other columns, hence the missing values will be imputed randomly. A day will be selected at random from the available days to fill the missing entries. This approach avoids making assumptions about the data while ensuring completeness in the dataset.

```

[383]: # List of days in a week
days_of_week = df['day_of_week'].dropna().unique()
days_of_week

# Replace missing values with a random day from the list

```

```
df['day_of_week'] = df['day_of_week'].apply(lambda x: random.  
↳choice(days_of_week) if pd.isnull(x) else x)
```


```
[384]: df['day_of_week'].value_counts()
```

```
[384]: day_of_week  
mon    1761  
thu    1673  
tue    1623  
wed    1619  
fri    1526  
Name: count, dtype: int64
```

4.3 Handling 999 values in pdays and pmonths

999 means that the customers are not previously contacted

```
[385]: # Calculate correlation matrix  
corr_matrix = df[num_cols].corr()  
  
# Create a heatmap  
plt.figure(figsize=(8, 6))  
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', linewidths=0.5, fmt='.2f')  
  
# Show the plot  
plt.title('Correlation Heatmap')  
plt.show()
```



‘nr.employed’, ‘emp.var.rate’, ‘euribor3m’ columns are highly correlated. These columns will be treated in Feature selection.

Based on the correlation matrix, it can be observed that the pdays and pmonths columns are highly correlated. Since these two features provide similar information, we will remove one of the columns to avoid redundancy and potential multicollinearity, which could affect model performance.

```
[386]: # Remove one of the correlated columns
df = df.drop(columns=['pmonths'])
```

To capture whether a customer has not been contacted previously, we will create a new column called previously_contacted


```
[387]: # Create a new column 'previously_contacted' based on 'pdays'
df['previously_contacted'] = df['pdays'].apply(lambda x: 0 if x == 999 else 1)
```

```
[388]: # plotting Bar Chart
plt.figure(figsize=(15,4))
```

```
ax = df['previously_contacted'].value_counts(normalize=True).plot(kind="bar")
plt.xticks(rotation= 0, fontsize=11)

# Add percentage labels to the top of bar
for p in ax.patches:
    ax.text(p.get_x()+p.get_width()/2., p.get_height()+0.002, f"{p.
    ↪get_height()*100:.1f}%", ha="center", fontsize=11)

plt.show()
```



The pdays column, which indicates the number of days since the client was last contacted, contains only 4% of the data marked as contacted previously. Given the sparsity of this information, it does not provide significant value for predictive modeling. Therefore, this column will be dropped to avoid unnecessary complexity in the model.

```
[389]: # Dropping pdays column due to sparsity of data
df = df.drop(columns=['pdays'])
```

5 Step 5: Feature Engineering

```
[390]: df[cat_cols].head()
```

```
[390]:
```

	profession	marital	schooling	default	housing	loan	contact	\
0	admin.	single	university.degree	no	no	yes	cellular	
1	services	single	high.school	no	no	no	cellular	
2	admin.	single	high.school	no	no	no	telephone	
3	admin.	divorced	university.degree	unknown	yes	no	cellular	
4	blue-collar	single	primary.education	unknown	yes	no	cellular	

	month	day_of_week	poutcome	responded
0	apr	wed	nonexistent	no
1	jul	thu	nonexistent	no
2	jun	tue	nonexistent	no
3	jul	tue	nonexistent	no


4 jul tue nonexistent no

5.1 Profession

```
[391]: # plotting Bar Chart
plt.figure(figsize=(15,4))
ax = df['profession'].value_counts(normalize=True).plot(kind="bar")
plt.xticks(rotation= 0, fontsize=11)

# Add percentage labels to the top of bar
for p in ax.patches:
    ax.text(p.get_x()+p.get_width()/2., p.get_height()+0.002, f"{p.
    ↪get_height()*100:.1f}%", ha="center", fontsize=11)

plt.show()
```



```
[392]: # Bivariate Analysis with Target column

# Get the order of categories based on value counts
profession_order = df['profession'].value_counts().index

plt.figure(figsize=(15,4))
ax = sns.countplot(x='profession', hue='responded', data=df,
    ↪order=profession_order)
plt.xticks(rotation=0, fontsize=11)

# Calculate total counts per 'profession' and 'responded' combination
total_counts = pd.crosstab(df['profession'], df['responded'])

# Add percentage labels to the top of each bar
for p in ax.patches:
    height = p.get_height()
    profession_name = ax.get_xticklabels()[round(p.get_x())].get_text() # Get
    ↪the profession name based on x position
```




```

# Calculate the total count for the current 'profession' and 'responded'
↪combination
total = total_counts.loc[profession_name, 'yes'] + total_counts.
↪loc[profession_name, 'no'] # Total for that profession

ax.text(p.get_x()+p.get_width()/2., p.get_height()+0.002, f"{height/
↪total*100:.1f}%", ha="center", fontsize=11)

plt.show()

```



Based on the findings from univariate and bivariate analysis,

- The top 5 professions account for 80% of the data, indicating that a significant portion of the dataset is concentrated in a few key professions.
- Retired individuals and students show distinct responses to the marketing campaign.

Based on these two observations,

1. Retired and students will be grouped into a new category called “Dependents” due to their distinct responses to the marketing campaign
2. Other less frequent professions will be combined into an “Others” category

```

[393]: # Create a mapping for profession categories
profession_mapping = {
    'retired': 'dependents',
    'student': 'dependents',
    'entrepreneur': 'others',
    'self-employed': 'others',
    'housemaid': 'others',
    'unemployed': 'others',
    'unknown': 'others'
}

# Apply the mapping to the 'profession' column
df['profession'] = df['profession'].map(profession_mapping).
↪fillna(df['profession'])


```

5.2 Marital

```
[394]: # plotting Bar Chart
plt.figure(figsize=(15,4))
ax = df['marital'].value_counts(normalize=True).plot(kind="bar")
plt.xticks(rotation= 0, fontsize=11)

# Add percentage labels to the top of bar
for p in ax.patches:
    ax.text(p.get_x()+p.get_width()/2., p.get_height()+0.002, f"{p.
    ↪get_height()*100:.1f}%", ha="center", fontsize=11)

plt.show()
```



The “unknown” category in the marital status column represents only 0.1% of the total records. Given its negligible size and the fact that it cannot be meaningfully grouped with other categories, these records are dropped to avoid introducing noise into the model.


```
[395]: # Drop records where marital status is 'unknown'
df = df[df['marital'] != 'unknown']
```

5.3 Schooling

```
[396]: # plotting Bar Chart
plt.figure(figsize=(10,3))
ax = df['schooling'].value_counts(normalize=True).plot(kind="bar")
plt.xticks(rotation= 0, fontsize=11)

# Add percentage labels to the top of bar
for p in ax.patches:
    ax.text(p.get_x()+p.get_width()/2., p.get_height()+0.002, f"{p.
    ↪get_height()*100:.1f}%", ha="center", fontsize=11)

plt.show()
```




5.4 Default

```
[397]: # plotting Bar Chart
plt.figure(figsize=(8,3))
ax = df['default'].value_counts(normalize=True).plot(kind="bar")
plt.xticks(rotation= 0, fontsize=11)

# Add percentage labels to the top of bar
for p in ax.patches:
    ax.text(p.get_x()+p.get_width()/2., p.get_height()+0.002, f"{p.
    ↪get_height()*100:.1f}%", ha="center", fontsize=11)

plt.show()
```



The “default” column contains only one record with a “yes” value, representing 0.01% of the data. To ensure meaningful analysis, this “yes” category is merged into the “unknown” category.

```
[398]: df['default'] = df['default'].replace('yes', 'unknown')
```

5.5 Housing

```
[399]: # Create subplots with 1 row and 2 columns
fig, axes = plt.subplots(1, 2, figsize=(14, 4))

# Bar chart for 'housing' column
ax1 = axes[0]
df['housing'].value_counts(normalize=True).plot(kind="bar", ax=ax1)
ax1.set_xticklabels(ax1.get_xticklabels(), rotation=0, fontsize=11)

# Add percentage labels to the top of bar
for p in ax1.patches:
    ax1.text(p.get_x() + p.get_width() / 2., p.get_height() + 0.002, f"{p.get_height() * 100:.1f}%", ha="center", fontsize=11)

# Bivariate Analysis with Target column
# Get the order of categories based on value counts
order = df['housing'].value_counts().index

ax2 = axes[1]
sns.countplot(x='housing', hue='responded', data=df, ax=ax2, order=order)
ax2.set_xticklabels(ax2.get_xticklabels(), rotation=0, fontsize=11)

# Calculate total counts per 'housing' and 'responded' combination
total_counts = pd.crosstab(df['housing'], df['responded'])


# Add percentage labels to the top of each bar
for p in ax2.patches:
    height = p.get_height()
    category_name = ax2.get_xticklabels()[round(p.get_x())].get_text() # Get
    ↪ the category name based on x position

    # Calculate the total count for the 'housing' and 'responded' combination
    total = total_counts.loc[category_name, 'yes'] + total_counts.
    ↪ loc[category_name, 'no'] # Total for that category

    ax2.text(p.get_x() + p.get_width() / 2., p.get_height() + 0.002, f"{height /
    ↪ total * 100:.1f}%", ha="center", fontsize=11)

# Adjust layout to avoid overlapping
plt.tight_layout()

# Show the plots
plt.show()
```




The “unknown” category in the “housing” column represents 2.2% of the data. Given its small proportion, it cannot be dropped without significant data loss, and it does not align with any other categories (yes, no). Therefore, the column will be left as is for analysis.

5.6 Loan

```
[400]: # plotting Bar Chart
plt.figure(figsize=(8,3))
ax = df['loan'].value_counts(normalize=True).plot(kind="bar")
plt.xticks(rotation= 0, fontsize=11)

# Add percentage labels to the top of bar
for p in ax.patches:
    ax.text(p.get_x()+p.get_width()/2., p.get_height()+0.002, f"{p.get_height()*100:.1f}%", ha="center", fontsize=11)

plt.show()
```




No changes have been made to the “loan” column.

5.7 Contact

```
[401]: # plotting Bar Chart
plt.figure(figsize=(8,3))
ax = df['contact'].value_counts(normalize=True).plot(kind="bar")
plt.xticks(rotation= 0, fontsize=11)

# Add percentage labels to the top of bar
for p in ax.patches:
    ax.text(p.get_x()+p.get_width()/2., p.get_height()+0.002, f"{p.
    ↪get_height()*100:.1f}%", ha="center", fontsize=11)

plt.show()
```




No changes have been made to the “contact” column.

5.8 Month

```
[402]: # plotting Bar Chart
plt.figure(figsize=(12,3))
ax = df['month'].value_counts(normalize=True).plot(kind="bar")
plt.xticks(rotation= 0, fontsize=11)

# Add percentage labels to the top of bar
for p in ax.patches:
    ax.text(p.get_x()+p.get_width()/2., p.get_height()+0.002, f"{p.
    ↪get_height()*100:.1f}%", ha="center", fontsize=11)

plt.show()
```



```
[403]: # Bivariate Analysis with Target column

# Get the order of categories based on value counts
order = df['month'].value_counts().index

plt.figure(figsize=(15,4))
ax = sns.countplot(x='month', hue='responded', data=df, order=order)
plt.xticks(rotation=0, fontsize=11)


# Calculate total counts per 'month' and 'responded' combination
total_counts = pd.crosstab(df['month'], df['responded'])

# Add percentage labels to the top of each bar
for p in ax.patches:
    height = p.get_height()
    category_name = ax.get_xticklabels()[round(p.get_x())].get_text() # Get
    ↳ the profession name based on x position

    # Calculate the total count for the current 'month' and 'responded'
    ↳ combination
    total = total_counts.loc[category_name, 'yes'] + total_counts.
    ↳ loc[category_name, 'no'] # Total for that profession

    ax.text(p.get_x()+p.get_width()/2., p.get_height()+0.002, f"{height/
    ↳ total*100:.1f}%", ha="center", fontsize=11)

plt.show()
```



- Based on the analysis, it is observed that most of the campaigns occur during the summer months(may, jun, jul, aug) along with nov.
- Additionally, campaigns conducted during less frequent months show a more balanced response rate, with a 50-50 split between “yes” and “no” responses.

```
[404]: # Grouping less frequent months into 'other' category

# Create a mapping for less frequent months
month_mapping = {
    'oct': 'others',
    'sep': 'others',
    'mar': 'others',
    'dec': 'others'
}


# Apply the mapping to the 'month' column
df['month'] = df['month'].map(month_mapping).fillna(df['month'])
```

5.9 Day of Week

```
[405]: # plotting Bar Chart
plt.figure(figsize=(12,3))
ax = df['day_of_week'].value_counts(normalize=True).plot(kind="bar")
plt.xticks(rotation= 0, fontsize=11)

# Add percentage labels to the top of bar
for p in ax.patches:
    ax.text(p.get_x()+p.get_width()/2., p.get_height()+0.002, f"{p.get_height()*100:.1f}%", ha="center", fontsize=11)

plt.show()
```


No changes have been made to the “day_of_week” column.

5.10 poutcome

```
[406]: # plotting Bar Chart
plt.figure(figsize=(12,3))
ax = df['poutcome'].value_counts(normalize=True).plot(kind="bar")
plt.xticks(rotation= 0, fontsize=11)

# Add percentage labels to the top of bar
for p in ax.patches:
    ax.text(p.get_x()+p.get_width()/2., p.get_height()+0.002, f"{p.
    ↪get_height()*100:.1f}%", ha="center", fontsize=11)

plt.show()
```



6 Step 6: Dealing with Imbalanced Data and Feature Selection

6.1 Dealing with Skewed Data in Numerical columns

Based on the analysis done on numerical columns, it is observed that custAge, campaign, previous and pastEmail columns are right skewed and nr.employed column is left skewed.





6.1.1 Transformation of Right Skewed Data





```
[407]: right_skew_col = ['custAge', 'campaign', 'previous', 'pastEmail']

pt = PowerTransformer(method='box-cox')

for col in right_skew_col:
    df[col+'_box'] = pt.fit_transform((df[col] + 1).values.reshape(-1, 1))
```

```
[408]: # Checking the skewness of transformed columns
for col in right_skew_col:
    fig, axes = plt.subplots(nrows=1, ncols=4, figsize=(15, 3))
    # Skewness and Histogram of transformed columns
    skewness_box = round(df[col+'_box'].skew(), 2)
    axes[0].text(0.5, 0.5, (f"{col+'_box'}\nSkewness: {skewness_box}"),
    ↪ fontsize=12, ha='center', va='center')
    axes[0].axis('off')
    axes[1].hist(df[col+'_box'])
    # Skewness and Histogram of actual columns
    skewness = round(df[col].skew(), 2)
    axes[2].text(0.5, 0.5, (f"{col}\nSkewness: {skewness}"), fontsize=12,
    ↪ ha='center', va='center')
    axes[2].axis('off')
    axes[3].hist(df[col])
    plt.show()
```





```
[409]: # Dropping the actual columns and renaming transformed columns
df = df.drop(columns = right_skew_col)

for col in right_skew_col:
    df = df.rename(columns={col+'_box': col})

df.head()
```

```
[409]:
```

	profession	marital	schooling	default	housing	loan	contact	\
0	admin.	single	university.degree	no	no	yes	cellular	
1	services	single	high.school	no	no	no	cellular	
2	admin.	single	high.school	no	no	no	telephone	
3	admin.	divorced	university.degree	unknown	yes	no	cellular	
4	blue-collar	single	primary.education	unknown	yes	no	cellular	

	month	day_of_week	poutcome	emp.var.rate	cons.price.idx	cons.conf.idx	\
0	apr	wed	nonexistent	-1.8	93.075	-47.1	
1	jul	thu	nonexistent	1.4	93.918	-42.7	
2	jun	tue	nonexistent	1.4	94.465	-41.8	
3	jul	tue	nonexistent	1.4	93.918	-42.7	
4	jul	tue	nonexistent	1.4	93.918	-42.7	

	euribor3m	nr.employed	responded	previously_contacted	custAge	campaign	\
0	1.498	5099.1	no		0 -0.577980	0.256015	
1	4.968	5228.1	no		0 -0.999330	2.196303	
2	4.961	5228.1	no		0 0.063493	-1.046201	
3	4.962	5228.1	no		0 1.254595	0.256015	

4	4.961	5228.1	no	0	0.032622	1.568621
---	-------	--------	----	---	----------	----------

	previous	pastEmail
0	-0.409207	-0.376481
1	-0.409207	-0.376481
2	-0.409207	-0.376481
3	-0.409207	-0.376481
4	-0.409207	-0.376481

6.1.2 Transformation of Left Skewed Data

```
[410]: df['nr.employed_square'] = df['nr.employed'].apply(lambda x: x**2)
df['nr.employed_log'] = np.log(df['nr.employed'] + 1)
df['nr.employed_sqrt'] = np.sqrt(df['nr.employed'])

print(f"Skewness of nr.employed: {round(df['nr.employed'].skew(),2)}")
print(f"Skewness of nr.employed_square: {round(df['nr.employed_square'].
↪skew(),2)}")
print(f"Skewness of nr.employed_log: {round(df['nr.employed_log'].skew(),2)}")
print(f"Skewness of nr.employed_sqrt: {round(df['nr.employed_sqrt'].skew(),2)}")
```

```
Skewness of nr.employed: -1.02
Skewness of nr.employed_square: -1.0
Skewness of nr.employed_log: -1.04
Skewness of nr.employed_sqrt: -1.03
```

Despite applying transformations to address the left-skewed data, there was little change in skewness. Therefore, no transformations will be applied to the 'nr.employed' column.

```
[411]: # Dropping the transformed columns
df = df.drop(columns = ['nr.employed_square', 'nr.employed_sqrt', 'nr.
↪employed_log'])

df.head()
```

```
[411]:
```

	profession	marital	schooling	default	housing	loan	contact	\
0	admin.	single	university.degree	no	no	yes	cellular	
1	services	single	high.school	no	no	no	cellular	
2	admin.	single	high.school	no	no	no	telephone	
3	admin.	divorced	university.degree	unknown	yes	no	cellular	
4	blue-collar	single	primary.education	unknown	yes	no	cellular	

	month	day_of_week	poutcome	emp.var.rate	cons.price.idx	cons.conf.idx	\
0	apr	wed	nonexistent	-1.8	93.075	-47.1	
1	jul	thu	nonexistent	1.4	93.918	-42.7	
2	jun	tue	nonexistent	1.4	94.465	-41.8	
3	jul	tue	nonexistent	1.4	93.918	-42.7	

4	jul	tue	nonexistent	1.4	93.918	-42.7
---	-----	-----	-------------	-----	--------	-------

	euribor3m	nr.employed	responded	previously_contacted	custAge	campaign \
0	1.498	5099.1	no	0	-0.577980	0.256015
1	4.968	5228.1	no	0	-0.999330	2.196303
2	4.961	5228.1	no	0	0.063493	-1.046201
3	4.962	5228.1	no	0	1.254595	0.256015
4	4.961	5228.1	no	0	0.032622	1.568621

	previous	pastEmail
0	-0.409207	-0.376481
1	-0.409207	-0.376481
2	-0.409207	-0.376481
3	-0.409207	-0.376481
4	-0.409207	-0.376481

6.2 Encoding and Standardization

```
[412]: # Redefining numerical and categorical columns
final_columns = df.columns
num_cols = df._get_numeric_data().columns
print(num_cols)

cat_cols = df.drop(columns=num_cols, axis=1).columns
# Dropping Target column
cat_cols = cat_cols.drop('responded')
print(cat_cols)
```

```
Index(['emp.var.rate', 'cons.price.idx', 'cons.conf.idx', 'euribor3m',
      'nr.employed', 'previously_contacted', 'custAge', 'campaign',
      'previous', 'pastEmail'],
      dtype='object')
Index(['profession', 'marital', 'schooling', 'default', 'housing', 'loan',
      'contact', 'month', 'day_of_week', 'poutcome'],
      dtype='object')
```

```
[413]: # Standardization of Numerical columns

# Initialize the StandardScaler
scaler = StandardScaler()

# Fit the scaler to the data and transform it
df[num_cols] = scaler.fit_transform(df[num_cols])
```

```
[414]: # One-Hot Encoding using pandas get_dummies
df = pd.get_dummies(df, columns=cat_cols, drop_first=True)
train_data_columns = df.drop(['responded'], axis=1).columns
```

```
train_data_columns
```

```
[414]: Index(['emp.var.rate', 'cons.price.idx', 'cons.conf.idx', 'euribor3m',
            'nr.employed', 'previously_contacted', 'custAge', 'campaign',
            'previous', 'pastEmail', 'profession_blue-collar',
            'profession_dependents', 'profession_management', 'profession_others',
            'profession_services', 'profession_technician', 'marital_married',
            'marital_single', 'schooling_primary.education',
            'schooling_professional.course', 'schooling_university.degree',
            'schooling_unknown', 'default_unknown', 'housing_unknown',
            'housing_yes', 'loan_unknown', 'loan_yes', 'contact_telephone',
            'month_aug', 'month_jul', 'month_jun', 'month_may', 'month_nov',
            'month_others', 'day_of_week_mon', 'day_of_week_thu', 'day_of_week_tue',
            'day_of_week_wed', 'poutcome_nonexistent', 'poutcome_success'],
          dtype='object')
```

6.3 Splitting X and y data

```
[415]: # Define X and y
X = df.drop(['responded'], axis=1)
print(X.shape)
y = df['responded']
y = y.map(dict(yes=1, no=0))
print(y.shape)
```

```
(8192, 40)
```

```
(8192,)
```


6.4 Feature Selection

```
[416]: # While performing EDA, we observed that these columns are highly correlated
corr_check = ['nr.employed', 'emp.var.rate', 'euribor3m']

# Compute the correlation matrix
corr_matrix = round(X[corr_check].corr(), 2)

# Plot the heatmap
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm')
```

```
[416]: <Axes: >
```



```
[417]: # Apply PCA
pca = PCA() # This will compute all principal components
X_pca = pca.fit_transform(X)

# Explained variance ratio for each component
explained_variance_ratio = pca.explained_variance_ratio_
print("Explained variance ratio per component:", explained_variance_ratio)

# Cumulative explained variance to decide how many components to keep
cumulative_variance = explained_variance_ratio.cumsum()
print("Cumulative explained variance:", cumulative_variance)
```

```
Explained variance ratio per component: [3.19643736e-01 1.11383998e-01
8.33919972e-02 7.21718034e-02
6.62848688e-02 5.19868492e-02 3.83186844e-02 2.90419449e-02
2.11449553e-02 1.79961370e-02 1.77451847e-02 1.73043166e-02
1.48933680e-02 1.42594454e-02 1.40817531e-02 1.26278015e-02
1.18204083e-02 1.00891079e-02 9.27437777e-03 8.29802748e-03
8.07255289e-03 6.86805707e-03 6.22102720e-03 5.77232060e-03
4.40075191e-03 4.24964958e-03 3.56111695e-03 3.27850640e-03
3.20573620e-03 2.96559473e-03 2.68833512e-03 2.61184280e-03
1.79194080e-03 1.32582313e-03 5.59058329e-04 2.97892750e-04]
```

```

2.41195887e-04 1.29826371e-04 5.97414438e-09 0.00000000e+00]
Cumulative explained variance: [0.31964374 0.43102773 0.51441973 0.58659154
0.6528764 0.70486325
0.74318194 0.77222388 0.79336884 0.81136497 0.82911016 0.84641448
0.86130784 0.87556729 0.88964904 0.90227684 0.91409725 0.92418636
0.93346074 0.94175877 0.94983132 0.95669938 0.9629204 0.96869272
0.97309347 0.97734312 0.98090424 0.98418275 0.98738848 0.99035408
0.99304241 0.99565426 0.9974462 0.99877202 0.99933108 0.99962897
0.99987017 0.99999999 1. 1. ]

```


```
[418]: pca.explained_variance_ratio_[:22].sum()
```

```
[418]: 0.9566993752712953
```

```

[419]: # Plot the cumulative explained variance
plt.plot(range(1, len(cumulative_variance)+1), cumulative_variance)
plt.xlabel('Number of Components')
plt.ylabel('Cumulative Explained Variance')
plt.title('Explained Variance vs. Number of Components')
plt.show()

```




```
[420]: # Choose the number of components to keep (e.g., 10)
pca = PCA(n_components=22)
X_pca = pca.fit_transform(X)

# Create a DataFrame with the reduced data
X = pd.DataFrame(X_pca, columns=[f'PC{i+1}' for i in range(22)])

print("Reduced Data (first few rows):")
print(X.head())
```

```
Reduced Data (first few rows):
```

	PC1	PC2	PC3	PC4	PC5	PC6	PC7	\
0	-1.591293	-1.968252	-0.910516	0.064336	0.282494	0.391136	0.545386	
1	1.677728	0.059129	-2.046948	0.466039	1.481641	-0.257416	0.608301	
2	1.865652	0.506196	-0.348914	-0.657122	-1.170769	0.810331	-0.246856	
3	1.602637	0.346158	0.186214	1.055461	-0.507925	-0.039870	0.903201	
4	1.703940	0.188707	-1.183672	1.002038	0.651079	-0.057463	0.415838	

	PC8	PC9	PC10	PC11	PC12	PC13	PC14	\
0	-0.717705	-0.264385	-0.336753	-0.488801	0.450286	-0.287174	-0.457923	
1	-0.388964	-0.430014	-0.026308	-0.533360	-0.100241	-0.562018	0.570689	
2	-0.915570	-0.585045	-0.281045	-0.518831	-0.223152	-0.184633	-0.410094	
3	-0.750353	-0.457941	0.287612	0.073004	0.506591	-0.255381	-0.583290	
4	0.492991	-1.374172	0.813999	-0.186786	-0.044126	-0.252455	-0.527293	

	PC15	PC16	PC17	PC18	PC19	PC20	PC21	\
0	0.805122	-0.197256	0.163059	0.020382	0.842613	-0.224935	0.007356	
1	-0.077696	0.604221	-0.616289	-0.210788	-0.036466	0.885973	0.165180	
2	-0.637562	-0.044971	0.801036	-0.086708	-0.027445	0.152358	0.307865	
3	-0.684990	-0.475250	-0.605413	0.660488	-0.090734	0.155662	-0.415321	
4	-0.694473	-0.225155	-0.450154	0.529223	-0.082392	0.031833	-0.351459	

	PC22
0	0.054657
1	0.315069
2	-0.008741
3	0.010436
4	-0.096269

6.0.1 Handling Imbalanced Data

The dataset used is highly imbalanced, with 88% of samples belonging to the “No” class and only 12% to the “Yes” class.

Oversampling the minority class may result in excessive duplication of data, leading to potential overfitting.

Undersampling the majority class could lead to a loss of valuable information.

To address this, we are proceeding with **mixed sampling** using SMOTE

```
[421]: # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
    ↪random_state=42)

# Apply SMOTE-NN for balancing the data
smote = SMOTE(random_state=42)

# Resample the training data
X_train_smote, y_train_smote = smote.fit_resample(X_train, y_train)

# Check the balance after resampling
print("Before resampling:")
print("Class distribution in training data:", np.bincount(y_train))
print("\nAfter resampling:")
print("Class distribution in resampled training data:", np.
    ↪bincount(y_train_smote))
```

Before resampling:

Class distribution in training data: [5066 668]

After resampling:

Class distribution in resampled training data: [5066 5066]

7 Step 7: Model Training

Model Selection

Binary classification tasks demand models that can effectively handle the nuances of the problem.

- Logistic Regression offers a straightforward and interpretable approach, particularly effective for linearly separable data.
- Models like Random Forest are advantageous for their ability to manage non-linear relationships and their robustness against overfitting.
- AdaBoost, on the other hand, uses boosting to iteratively improve weak learners, making it effective at handling complex patterns, particularly in noisy data.
- Ridge Regression, with its L2 regularization, helps prevent overfitting in high-dimensional datasets, especially when linear relationships are present.

Metric Selection

In evaluating the performance of a binary classification model, **accuracy** is often considered a primary metric. However, since the data is imbalanced, focusing solely on accuracy can be misleading. Metrics such as **precision** and **recall** provide a more nuanced understanding by addressing the costs associated with false positives and false negatives. Additional metrics like ROC-AUC are instrumental in offering a balanced and comprehensive evaluation.

7.1 Defining ML models

```
[422]: # Logistic Regression model
logreg = LogisticRegression(max_iter=1000)
logreg.fit(X_train_smote, y_train_smote)
y_pred_logreg = logreg.predict(X_test)

# Random Forest Classifier model
rf_model = RandomForestClassifier(class_weight='balanced', n_estimators=100,
    random_state=42)
rf_model.fit(X_train_smote, y_train_smote)
y_pred_rf = rf_model.predict(X_test)

# Ridge Classifier model
ridge = RidgeClassifier()
ridge.fit(X_train_smote, y_train_smote)
y_pred_ridge = ridge.predict(X_test)

# AdaBoost Classifier model
adaboost = AdaBoostClassifier(n_estimators=50, random_state=42)
adaboost.fit(X_train_smote, y_train_smote)
y_pred_adaboost = adaboost.predict(X_test)
```

7.2 Calculating Metrics

```
[423]: # Calculating metrics for each model

# Defining all models in a list
model = {
    "Logistic Regression": y_pred_logreg,
    "Ridge Classifier": y_pred_ridge,
    "AdaBoost Classifier": y_pred_adaboost,
    "Random Forest Classifier": y_pred_rf
}

# Metrics calculation
for name, y_pred in model.items():
    print(f"\n===== {name} =====\n")
    # Confusion matrix
    conf_matrix = confusion_matrix(y_test, y_pred)
    print(f"Confusion Matrix:\n{conf_matrix}\n")

    # Accuracy
    print(f"\nAccuracy: {round(accuracy_score(y_test, y_pred), 2)}\n")

    # Classification report
    clf_report = classification_report(y_test, y_pred)
    print(f"\nClassification Report:\n{clf_report}\n")
```

```
# ROC-AUC Score
roc_auc = roc_auc_score(y_test, y_pred)
print(f"ROC-AUC Score: {round(roc_auc,2)}\n")
```

===== Logistic Regression =====

Confusion Matrix:

```
[[1741  459]
 [  91 167]]
```

Accuracy: 0.78

Classification Report:

	precision	recall	f1-score	support
0	0.95	0.79	0.86	2200
1	0.27	0.65	0.38	258
accuracy			0.78	2458
macro avg	0.61	0.72	0.62	2458
weighted avg	0.88	0.78	0.81	2458

ROC-AUC Score: 0.72

===== Ridge Classifier =====

Confusion Matrix:

```
[[1715  485]
 [  86 172]]
```

Accuracy: 0.77

Classification Report:

	precision	recall	f1-score	support
0	0.95	0.78	0.86	2200
1	0.26	0.67	0.38	258
accuracy			0.77	2458

macro avg	0.61	0.72	0.62	2458
weighted avg	0.88	0.77	0.81	2458

ROC-AUC Score: 0.72

===== AdaBoost Classifier =====

Confusion Matrix:

```
[[1724  476]
 [  91 167]]
```

Accuracy: 0.77

Classification Report:

	precision	recall	f1-score	support
0	0.95	0.78	0.86	2200
1	0.26	0.65	0.37	258
accuracy			0.77	2458
macro avg	0.60	0.72	0.61	2458
weighted avg	0.88	0.77	0.81	2458

ROC-AUC Score: 0.72

===== Random Forest Classifier =====

Confusion Matrix:

```
[[2022  178]
 [ 145 113]]
```

Accuracy: 0.87

Classification Report:

	precision	recall	f1-score	support
0	0.93	0.92	0.93	2200
1	0.39	0.44	0.41	258
accuracy			0.87	2458

macro avg	0.66	0.68	0.67	2458
weighted avg	0.88	0.87	0.87	2458

ROC-AUC Score: 0.68

Random Forest Classifier, with its strong performance in accuracy, F1-score, and recall, is the best model for predicting customer responses. Its high accuracy ensures reliable predictions, while the F1-score reflects a good balance between precision and recall, especially for identifying the minority class.

7.3 Hyperparameter Tuning of Random Forest Classifier

```
[ ]: from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import randint

# Define the Random Forest model
rf = RandomForestClassifier(random_state=42)

# Define the hyperparameters distribution to sample from
param_dist = {
    'n_estimators': randint(50, 200), # Number of trees
    'max_depth': [None, 10, 20, 30, 40], # Maximum depth of the tree
    'min_samples_split': randint(2, 20), # Minimum number of samples required
    ↪to split a node
    'min_samples_leaf': randint(1, 20), # Minimum number of samples required
    ↪at a leaf node
    'max_features': ['auto', 'sqrt', 'log2'], # The number of features to
    ↪consider for the best split
    'bootstrap': [True, False] # Whether to use bootstrap samples
}

# Set up RandomizedSearchCV
random_search = RandomizedSearchCV(estimator=rf, param_distributions=param_dist,
                                   n_iter=100, cv=5, scoring='accuracy',
                                   ↪n_jobs=-1, verbose=2, random_state=42)

# Fit the random search
random_search.fit(X_train_smote, y_train_smote)

# Print the best hyperparameters
print(f"Best hyperparameters: {random_search.best_params_}")

# Evaluate the model with the best hyperparameters
best_rf_random = random_search.best_estimator_
y_pred_rf_random = best_rf_random.predict(X_test)
```

```
# Print classification report
print(classification_report(y_test, y_pred_rf_random))
```

Fitting 5 folds for each of 100 candidates, totalling 500 fits
 Best hyperparameters: {'bootstrap': False, 'max_depth': None, 'max_features': 'log2', 'min_samples_leaf': 1, 'min_samples_split': 6, 'n_estimators': 110}

	precision	recall	f1-score	support
0	0.93	0.93	0.93	2200
1	0.39	0.38	0.39	258
accuracy			0.87	2458
macro avg	0.66	0.66	0.66	2458
weighted avg	0.87	0.87	0.87	2458

```
[424]: # Best hyperparameters for Random Forest Classifier using RandomSearchCV
best_params = {
    'bootstrap': False,
    'max_features': 'log2',
    'min_samples_leaf': 1,
    'min_samples_split': 6,
    'n_estimators': 110
}

# Create the RandomForestClassifier with the best hyperparameters
final_model = RandomForestClassifier(**best_params)

# Fit the model to the training data
final_model.fit(X_train, y_train)

# Make predictions on the test data
y_pred = final_model.predict(X_test)
```

```
[425]: # Confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)
print(f"Confusion Matrix:\n{conf_matrix}\n")

# Accuracy
print(f"\nAccuracy: {round(accuracy_score(y_test, y_pred),2)}\n")

# Classification report
clf_report = classification_report(y_test, y_pred)
print(f"\nClassification Report:\n{clf_report}\n")

# ROC-AUC Score
roc_auc = roc_auc_score(y_test, y_pred)
```

```
print(f"ROC-AUC Score: {round(roc_auc,2)}\n")
```

Confusion Matrix:

```
[[2127   73]
 [ 176   82]]
```

Accuracy: 0.9

Classification Report:

	precision	recall	f1-score	support
0	0.92	0.97	0.94	2200
1	0.53	0.32	0.40	258
accuracy			0.90	2458
macro avg	0.73	0.64	0.67	2458
weighted avg	0.88	0.90	0.89	2458

ROC-AUC Score: 0.64

8 Step 8: Model Deployment

8.1 Defining function for Data Cleaning

```
[426]: def data_cleaning(test):

    # Defining 'previously_contacted' column based on 'pdays'
    test['previously_contacted'] = test['pdays'].apply(lambda x: 0 if x == 999
    ↪ else 1)

    # Selecting the final columns
    final_columns = ['profession', 'marital', 'schooling', 'default', 'housing',
    ↪ 'loan', 'contact', 'month', 'day_of_week', 'poutcome', 'emp.var.rate',
    ↪ 'cons.price.idx', 'cons.conf.idx', 'euribor3m', 'nr.employed',
    ↪ 'previously_contacted', 'custAge', 'campaign', 'previous', 'pastEmail']
    test = test[final_columns]

    # Replace missing values
    # Schooling - data aggregation
    test['schooling'] = test['schooling'].replace(['basic.4y', 'basic.6y', 'basic.
    ↪ 9y'], 'primary.education')
    test['schooling'] = test['schooling'].replace('illiterate', 'unknown')
```



```

# Replacing missing data using profession column
schooling_profession_mapping = {
    'technician': 'professional.course',
    'blue-collar': 'primary.education',
    'admin.': 'university.degree',
    'services': 'high.school'}

# Function to impute missing 'Schooling' values based on 'Profession'
def impute_schooling(row):
    if pd.isnull(row['schooling']):
        return schooling_profession_mapping.get(row['profession'], 'unknown')
    else:
        return row['schooling']

# Apply the function to impute missing values in 'Schooling'
test['schooling'] = test.apply(impute_schooling, axis=1)

# Imputing missing values for 'custAge'
test['custAge'] = test.apply(
    lambda row: mean_ages['retired'] if row['profession'] == 'retired' and pd.
↪isnull(row['custAge']) else
        mean_ages['student'] if row['profession'] == 'student' and pd.
↪isnull(row['custAge']) else
        mean_ages['other'] if pd.isnull(row['custAge']) else
↪row['custAge'],
    axis=1
)

# Imputing missing values for Day of week
test['day_of_week'] = test['day_of_week'].apply(lambda x: random.
↪choice(days_of_week) if pd.isnull(x) else x)

return test

```

8.2 Defining function for Data Preprocessing

```

[427]: def preprocess_data(test, scaler=scaler, pt=pt, pca=pca, train_data_columns=
↪train_data_columns, num_cols=num_cols , cat_cols=cat_cols):
    """
    Apply the same transformations to the test data as done for the training
    ↪data.

    Parameters:
    - test (DataFrame): The test data to preprocess.
    - scaler (StandardScaler): The fitted StandardScaler from training data.
    - pt (PowerTransformer): The fitted PowerTransformer from training data.
    - pca (PCA): The fitted PCA transformer from training data.

```

```

- train_data_columns (list): The list of columns from the training data.
- num_cols (list): The list of numerical columns.
- cat_cols (list): The list of categorical columns.

Returns:
- test (DataFrame): The preprocessed and PCA-transformed test data.
"""

# 1. Apply profession mapping (using predefined mapping)
test['profession'] = test['profession'].map({
    'retired': 'dependents',
    'student': 'dependents',
    'entrepreneur': 'others',
    'self-employed': 'others',
    'housemaid': 'others',
    'unemployed': 'others',
    'unknown': 'others'
}).fillna(test['profession'])

# 2. Replace 'yes' with 'unknown' in 'default' column
test['default'] = test['default'].replace('yes', 'unknown')

# 3. Apply month mapping (using predefined mapping)
test['month'] = test['month'].map({
    'oct': 'others',
    'sep': 'others',
    'mar': 'others',
    'dec': 'others'
}).fillna(test['month'])

# 4. Apply Box-Cox transformation to numerical columns using fitted
↳ PowerTransformer
right_skew_col = ['custAge', 'campaign', 'previous', 'pastEmail']

for col in right_skew_col:
    test[col] = pt.transform((test[col] + 1).values.reshape(-1, 1))

# 5. Standardize numerical columns using the fitted scaler
test[num_cols] = scaler.transform(test[num_cols])

# 6. Apply one-hot encoding to categorical columns
test = pd.get_dummies(test, columns=cat_cols, drop_first=True)

# 7. Ensure the test data has the same columns as the training data (in
↳ case of missing categories)
test = test.reindex(columns=train_data_columns, fill_value=0)

```

```

# 8. Apply PCA transformation to the test data
# Transform the test data with PCA (based on the number of components
↳fitted during training)
test_pca = pca.transform(test)

# Convert PCA result to a DataFrame with column names like 'PC1', 'PC2', ...
↳, 'PCn'
test = pd.DataFrame(test_pca, columns=[f'PC{i+1}' for i in range(test_pca.
↳shape[1])])

# Return the preprocessed test data after PCA transformation
return test

```

8.3 Creating Pipeline

```

[428]: # Define the pipeline for data cleaning and preprocessing
data_pipeline = Pipeline([
    ('data_cleaning', FunctionTransformer(func=data_cleaning)),
    ('data_preprocessing', FunctionTransformer(func=preprocess_data)),
])

# Save the data pipeline to disk
joblib.dump(data_pipeline, 'data_pipeline.joblib')

# Save the trained model to disk
joblib.dump(final_model, 'trained_model.joblib')

```

```

[428]: ['trained_model.joblib']

```

8.4 Loading Test Data and Predicting

```

[429]: test_data = pd.read_excel("/content/drive/MyDrive/Upgrad/Data sets/Capstone/
↳test.xlsx")

```

```

[430]: # Load the trained model and data pipeline from disk
trained_model = joblib.load('trained_model.joblib')
data_pipeline = joblib.load('data_pipeline.joblib')

```

```

[431]: # Apply the preprocessing steps to the test data
processed_test_data = data_pipeline.transform(test_data)

# Predict outcomes using the trained model on the processed test data
predictions = trained_model.predict(processed_test_data)

# Add the predictions as a new column in the processed test data
processed_test_data['Prediction'] = predictions

```

```
[433]: # Save the processed test data with predictions to an Excel file  
output_file = 'predictions_output.xlsx'  
processed_test_data.to_excel(output_file, index=False)
```